

# CSC453

## Parallel Processing Project

Ammar Alamri

Std#: 438104833

### Bitonic Parallel Sort Implementation:

Note: I tried my best to replace module (%), division and pow() operations, with shift and bitwise operations, to increase performance and provide a more unique solution, since this might affect readability I also added comments showing how the operations would be done using traditional methods.

```
// BitonicSorter works only on arrays of size 2^n.
__global__ void BitonicSorter(int* in, int N, int max_iterations) {
    int index = (threadIdx.x);
    int seq_length = 0, shift = 0;

    for (int i = 1; i < max_iterations + 1; i++) { // for steps, for stages.
        for (int j = 1; j < i + 1; j++) {

            // 1 << i-j+1 = pow(2, i-j+1).
            seq_length = 1 << i - j + 1;

            // seq_length / 2.
            shift = seq_length >> 1;

            // index % seq_length < shift.
            if ((index & (seq_length - 1)) < shift) {

                // if (index / pow(2,i)) is even.
                if ((index >> i & 1) == 0) {
                    if (in[index] > in[index + shift]) {
                        int temp = in[index];
                        in[index] = in[index + shift];
                        in[index + shift] = temp;
                    }
                }
                else if (in[index] < in[index + shift]) {
                    int temp = in[index];
                    in[index] = in[index + shift];
                    in[index + shift] = temp;
                }
            }
            __syncthreads();
        }
    }
}
```

```

int main() {
    /* =main function= */
    int N = 50;
    int fillers = 0, max_iterations = 0, altSize = 0;

    /* allocation and array generation */
    int* a, * b, * d_a;
    int size = sizeof(int) * N;

    a = (int*)malloc(size);
    b = (int*)malloc(size);

    cudaMalloc(&d_a, size);

    random_ints(a, N);

    // Checking if size is 2^n.
    for (int i = 0; i < 20; i++) {
        // altSize = pow(2, i)
        altSize = 1 << i;

        if (N <= altSize) {
            fillers = altSize - N;
            max_iterations = i;
            break;
        }
    }
    // if original size not 2^n, we add fillers to the end of the array.
    if (fillers > 0) {
        size = sizeof(int) * altSize;
        realloc(a, size);

        for (int i = 0; i < fillers; i++)
            a[N + i] = INT_MAX;
    }

    cudaMalloc(&d_a, size);
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);

    BitonicSorter << <1, altSize >> > (d_a, altSize, max_iterations);
    cudaDeviceSynchronize();

    cudaMemcpy(a, d_a, size, cudaMemcpyDeviceToHost);

    // Trim output to original size N.
    memcpy(b, a, sizeof(int) * N);
    printf("After sorting, array is: "); printArray(b, N);
    return 0;
}

/* Helper functions */
void random_ints(int*& array, int size) {
    for (int i = 0; i < size; i++)
        array[i] = rand() % size;
}

void printArray(int* array, int size) {
    for (int i = 0; i < size - 1; i++)
        printf("%d, ", array[i]);
    printf("%d.\n", array[size - 1]);
}

```